



SAVITRIBAI PHULE PUNE UNIVERSITY

Mobile Computing

COMPUTER SCIENCE
Fourth Year (2019 Pattern)

Unit-1 (NOTES)

PUNE
UNIVERSITY

UNIT I: Introduction To Machine Learning

1. Introduction to Machine Learning

Definition: Machine learning is a subfield of artificial intelligence (AI) that focuses on the development of algorithms and statistical models that enable computer systems to perform tasks by learning from data, without being explicitly programmed. Instead of following a set of predefined instructions, ML models use patterns and inferences from a given dataset to make predictions or decisions. The core idea is to allow computers to learn and improve automatically over time.

Detailed Explanation:

- **Learning from Data:** The fundamental principle of machine learning is to build systems that can learn from and make predictions based on data. These systems can identify patterns, relationships, and trends within the data without human intervention.
- **Algorithmic Foundation:** Machine learning relies on algorithms that parse data, learn from it, and then apply what they've learned to make an informed decision or prediction. These algorithms are the core of any ML model.
- **Iterative Process:** The process of training a machine learning model is often iterative. The model is trained on a dataset, its performance is evaluated, and then the model is fine-tuned to improve its accuracy and efficiency.
- **Generalization:** A key goal of machine learning is for the model to generalize well, meaning it should perform accurately on new, unseen data, not just the data it was trained on. This is what distinguishes a good model from one that has simply memorized the training data (overfitting).

Mathematical / Technical Aspect: At its core, many machine learning algorithms are based on optimization problems. For example, in linear regression, the goal is to find the line that best fits the data by minimizing the sum of squared errors. The formula for a simple linear regression model is: $y = \beta_0 + \beta_1 x + \epsilon$ Where:

- y is the dependent variable
- x is the independent variable
- β_0 is the y -intercept

- β_1 is the slope
- ϵ represents the error term

Real-World Examples:

1. **Email Spam Detection:** Gmail uses machine learning to identify and filter spam emails. The model is trained on a massive dataset of emails labeled as either "spam" or "not spam" to learn the characteristics of spam.
2. **Predictive Text:** The keyboard on your smartphone uses an ML model that analyzes your past writing habits to predict the next word you are likely to type.
3. **Netflix Recommendations:** Netflix's recommendation engine is a classic example of machine learning. It analyzes your viewing history and the viewing habits of millions of other users to suggest movies and shows you might enjoy.

Applications:

- **Healthcare:** Disease diagnosis, personalized medicine, and drug discovery.
- **Finance:** Fraud detection, algorithmic trading, and credit scoring.
- **Retail:** Personalized recommendations, inventory management, and demand forecasting.
- **Autonomous Vehicles:** Self-driving cars use machine learning for object detection, navigation, and decision-making.

Important Notes / Key Takeaways:

- Machine learning is a subset of AI.
 - The primary function is to learn from data, not from explicit programming.
 - A crucial concept is generalization, which ensures the model works on new data.
 - It's a data-driven approach, making the quality and quantity of data paramount.
-

2. Comparison of Machine Learning with Traditional Programming

Definition: Traditional programming involves a developer writing specific instructions for a computer to follow to achieve a particular task. In contrast, machine learning involves feeding a computer with data and an algorithm, allowing it to learn the rules or patterns from the data on its own. The core difference lies in who defines the rules: the programmer in traditional programming or the data in machine learning.

Detailed Explanation:

- **Traditional Programming:**

- **Input:** Data + Rules
- **Process:** The program follows a set of explicit instructions (rules) created by a programmer.
- **Output:** The result is determined by the fixed rules and input data.
- **Example:** A program to calculate the factorial of a number is based on a specific, non-changing set of instructions.

- **Machine Learning:**

- **Input:** Data + Algorithm
- **Process:** The algorithm processes the data to learn the rules or patterns.
- **Output:** The result (the "rules" or a trained model) can be applied to new data.
- **Example:** A spam filter learns to identify spam emails by analyzing a dataset of emails and their labels, not from a programmer's pre-defined rules.

Real-World Examples:

1. **Face Recognition:**

- **Traditional Programming:** You would write code with very specific rules, like "If the image has a nose, two eyes, and a mouth in these precise pixel locations, it's a face." This approach is rigid and fails easily with variations.
- **Machine Learning:** You provide thousands of images of faces and non-faces. The ML model learns the features that distinguish a face, such as the shape, spacing, and textures, on its own.

2. **Weather Prediction:**

- **Traditional Programming:** A program would use fixed equations of meteorology to calculate the weather based on input data like temperature, pressure, and humidity.
- **Machine Learning:** An ML model would be trained on historical weather data to find complex, non-linear patterns that predict future weather, often with higher accuracy.

3. **Language Translation:**

- **Traditional Programming:** A program would use a huge dictionary and grammatical rules to translate. This is often clunky and fails to capture context.

- **Machine Learning:** An ML model is trained on millions of sentences and their translations, learning the nuances and context of language to provide more natural-sounding translations.

Applications:

- **Traditional Programming:** Used for tasks that have well-defined, static rules, such as operating systems, compilers, and simple calculators.
- **Machine Learning:** Used for complex tasks where the rules are not easily defined, such as image recognition, natural language processing, and personalized recommendations.

Important Notes / Key Takeaways:

- Traditional programming is rule-based, while machine learning is data-driven.
 - Traditional programming is deterministic; the output is always the same for the same input. ML can be probabilistic or non-deterministic.
 - ML excels at problems that are too complex for a programmer to define all the rules for.
-

3. ML vs AI vs Data Science

Definition: This section clarifies the relationships between three often-confused fields: Artificial Intelligence (AI), Machine Learning (ML), and Data Science. In simple terms, AI is the broad concept of creating intelligent machines, ML is a specific method used to achieve AI, and Data Science is an interdisciplinary field that uses various techniques (including ML) to extract knowledge and insights from data.

Detailed Explanation:

- **Artificial Intelligence (AI):**
 - **Concept:** The overarching science of making machines intelligent.
 - **Goal:** To create systems that can reason, learn, perceive, and act autonomously, mimicking human cognitive functions.
 - **Scope:** Very broad, includes fields like robotics, natural language processing, computer vision, and expert systems.
- **Machine Learning (ML):**

- **Concept:** A subset of AI that uses algorithms and statistical models to enable machines to learn from data without explicit programming.
- **Goal:** To enable computers to find patterns in data and make predictions.
- **Relationship:** ML is one of the primary methods used to build AI systems.
- **Data Science:**
 - **Concept:** An interdisciplinary field that uses scientific methods, processes, algorithms, and systems to extract knowledge and insights from structured and unstructured data.
 - **Goal:** To gain insights and understanding from data.
 - **Scope:** Broader than ML; it involves data collection, cleaning, analysis, visualization, and communication of results. It uses techniques from statistics, programming, and mathematics. Machine learning is often a tool used by data scientists for predictive analysis.

Real-World Examples:

1. A Self-Driving Car:

- **AI:** The entire system of the car, which includes navigation, object detection, and decision-making, is an AI application.
- **ML:** The car's object detection system, which identifies pedestrians and other cars, is powered by a machine learning model.
- **Data Science:** The analysis of traffic data to optimize routes and the analysis of driving habits to improve safety features is part of data science.

2. Fraud Detection in Banking:

- **AI:** The entire system designed to flag suspicious transactions is an AI system.
- **ML:** The specific algorithm that learns from historical transaction data to identify new fraudulent patterns is a machine learning model.
- **Data Science:** The process of a data scientist collecting transaction data, cleaning it, analyzing it to identify trends, and creating dashboards to monitor fraud is a data science task.

3. E-commerce Recommendation Systems:

- **AI:** The "smart" shopping experience, including personalized product displays and search results, is an AI application.
- **ML:** The collaborative filtering or content-based recommendation models that suggest products are machine learning models.

- **Data Science:** The data scientist's role would be to analyze customer behavior data, identify purchasing patterns, and use the insights to improve the recommendation system's effectiveness.

Applications:

- **AI:** Used in robotics, expert systems, and natural language processing to mimic human intelligence.
- **ML:** Used for tasks like image recognition, predictive modeling, and speech recognition.
- **Data Science:** Used to generate insights for business strategy, scientific research, and policy-making.

Important Notes / Key Takeaways:

- AI is the big picture, ML is a specific technique within that picture, and Data Science is the field that uses these techniques to derive insights from data.
 - The relationship is often described as concentric circles: ML is a subset of AI, and both are often used within the broader field of Data Science.
-

4. Types of Learning

Definition: Machine learning models are typically categorized based on the nature of the training data and the type of problem they are designed to solve. The main types of learning are supervised, unsupervised, semi-supervised, and reinforcement learning, each with a distinct approach to problem-solving.

4.1 Supervised Learning

Definition: Supervised learning is a type of machine learning where an algorithm learns from a labeled dataset. The dataset contains input features (the independent variables) and their corresponding correct output values (the dependent variable). The goal of the algorithm is to learn a mapping function from the input to the output so that it can accurately predict the output for new, unseen data.

Detailed Explanation:

- **Training with Labels:** The key characteristic is the use of labeled data. The model is essentially "supervised" by a human or a pre-existing system that has provided the correct answers.
- **Two Main Tasks:** Supervised learning problems are generally divided into two types:
 - **Regression:** The output is a continuous numerical value (e.g., predicting house prices, temperature).
 - **Classification:** The output is a discrete category or class (e.g., classifying an email as spam or not spam, identifying an image as a cat or a dog).
- **Generalization:** The model learns a relationship between the input features and the output labels. The more diverse and representative the training data, the better the model's ability to generalize to new data.

Mathematical / Technical Aspect: In regression, a simple linear model might be used: $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$ Where:

- $\hat{y}$ is the predicted output
- $x_1, x_2, \dots, x_n$ are the input features
- $\beta_0, \beta_1, \dots, \beta_n$ are the coefficients (parameters) learned by the model

Real-World Examples:

1. **House Price Prediction:** A model is trained on a dataset of houses with features like square footage, number of bedrooms, and location, and their corresponding sale prices. It learns to predict the price of a new house.
2. **Image Classification:** A model is trained on a large dataset of images, each labeled with the object it contains (e.g., "car," "truck," "bus"). It can then classify new images.
3. **Credit Risk Assessment:** Banks use supervised learning models trained on historical customer data (income, debt, credit history) to classify new applicants as low-risk or high-risk for a loan.

Applications:

- **Medical Diagnosis:** Predicting a disease based on patient symptoms.
- **Market Forecasting:** Predicting stock prices or sales figures.
- **Sentiment Analysis:** Classifying a review as positive, negative, or neutral.

Important Notes / Key Takeaways:

- Requires a labeled dataset for training.
- The model learns a mapping from inputs to outputs.
- Used for both regression and classification problems.
- Performance is highly dependent on the quality of the labels.

4.2 Unsupervised Learning

Definition: Unsupervised learning is a machine learning technique where the algorithm is given an unlabeled dataset and must find patterns, structures, or relationships within the data on its own. The model is not given any correct answers or outputs to guide its learning. The primary goal is to explore the data and discover hidden insights.

Detailed Explanation:

- **No Labels:** The key difference from supervised learning is the absence of a target variable or output label. The model is left to find its own structure.
- **Two Main Tasks:** Unsupervised learning problems are typically used for:
 - **Clustering:** Grouping similar data points together.
 - **Association:** Discovering rules that describe large portions of the data.
- **Discovering Hidden Structure:** The algorithm's objective is to find meaningful insights that might not be obvious to a human analyst. It identifies natural groupings or patterns.

Mathematical / Technical Aspect: In clustering, a common algorithm like K-Means aims to minimize the within-cluster sum of squares (WCSS), which can be expressed as: $WCSS = \sum_{i=1}^k \sum_{x \in S_i} ||x - \mu_i||^2$ Where:

- S_i is the set of data points in cluster i .
- μ_i is the centroid of cluster i .
- k is the number of clusters.

Real-World Examples:

1. **Customer Segmentation:** A marketing team uses clustering to group customers based on their purchasing behavior, demographics, and interests to create targeted marketing campaigns.
2. **Anomaly Detection:** Identifying fraudulent transactions in a massive dataset by detecting transactions that don't fit the normal patterns.

3. **Market Basket Analysis:** E-commerce sites use association rules to discover which products are frequently bought together, leading to "Customers who bought this also bought..." recommendations.

Applications:

- **Genetics:** Grouping gene expressions with similar functionality.
- **Cybersecurity:** Detecting unusual network traffic patterns that could indicate a cyberattack.
- **Social Media Analysis:** Grouping users with similar interests to form communities.

Important Notes / Key Takeaways:

- Does not require a labeled dataset.
- The model finds patterns and structures on its own.
- Primarily used for clustering, association, and dimensionality reduction.

4.3 Semi-Supervised Learning

Definition: Semi-supervised learning is a machine learning approach that combines a small amount of labeled data with a large amount of unlabeled data during training. The goal is to leverage the vast amount of unlabeled data to improve the learning process and achieve better performance than supervised learning alone, especially when obtaining labeled data is expensive or time-consuming.

Detailed Explanation:

- **Combination of Labeled and Unlabeled Data:** This approach is a hybrid of supervised and unsupervised learning. The model first learns from the labeled data and then uses that knowledge to infer labels for the unlabeled data, which it then uses for further training.
- **Cost-Effectiveness:** It addresses the problem of data scarcity. In many real-world scenarios, it is easy to collect vast amounts of raw data (unlabeled), but very difficult and expensive to label it manually.
- **Improved Performance:** By using both types of data, the model can learn more robust and generalizable features, leading to higher accuracy compared to a model trained only on the small labeled dataset.

Real-World Examples:

1. **Web Content Classification:** A search engine has a small number of web pages manually labeled (e.g., "sports," "news") but millions of unlabeled pages. A semi-supervised model can use the labeled data to infer the categories of the unlabeled pages, improving its classification accuracy.
2. **Speech Recognition:** A model is trained with a small dataset of labeled audio files (speech and text transcriptions) and a vast amount of unlabeled audio. The model learns from the labeled data and uses the unlabeled data to better understand the nuances of spoken language.
3. **Medical Image Segmentation:** A small number of medical scans are manually labeled by experts to identify tumors. A semi-supervised model can then be trained on these and a large number of unlabeled scans to improve its ability to segment tumors in new images.

Applications:

- **Face Recognition:** Training a model with a few labeled photos and thousands of unlabeled ones.
- **Natural Language Processing:** For tasks like part-of-speech tagging or sentiment analysis when labeled data is limited.
- **Text Document Classification:** Categorizing large document corpora.

Important Notes / Key Takeaways:

- A blend of supervised and unsupervised learning.
- Addresses the challenge of limited labeled data.
- Often leads to better performance than using labeled data alone.

4.4 Reinforcement Learning

Definition: Reinforcement learning (RL) is a machine learning paradigm concerned with how an "agent" should take actions in an "environment" to maximize a cumulative "reward." The agent learns through trial and error, receiving rewards for good actions and penalties for bad ones. It doesn't use a labeled dataset but learns from the consequences of its actions.

Detailed Explanation:

- **Agent, Environment, and Reward:**
 - **Agent:** The learner or decision-maker.

- **Environment:** The world the agent interacts with.
- **Reward:** A feedback signal that indicates how well the agent performed.
- **Learning through Interaction:** The agent performs an action, the environment changes state, and the agent receives a reward or penalty. The agent's goal is to learn a policy (a strategy) that maps states to actions to maximize its total reward over time.
- **No Labeled Data:** Unlike supervised learning, there's no "correct" answer provided. The learning is based on the feedback from the environment, which is a key differentiator.

Mathematical / Technical Aspect: The core of many reinforcement learning problems is the Bellman equation, which relates the value of a state to the values of its successor states: $V(s) = R(s,a) + \gamma \cdot \max_a V(s')$ Where:

- $V(s)$ is the value of state s .
- $R(s,a)$ is the immediate reward for taking action a in state s .
- γ (gamma) is the discount factor, which determines the importance of future rewards.

Real-World Examples:

1. **Game Playing:** An AI agent plays chess or Go against a human or itself. It receives a reward for winning and a penalty for losing, learning the best moves through countless games.
2. **Robotics:** A robot learns to walk by being given a positive reward for moving forward and a penalty for falling down. It learns the motor controls to maintain balance.
3. **Resource Management:** An RL agent can be trained to optimize energy usage in a data center by taking actions (e.g., turning on/off servers) to minimize energy consumption while meeting computational demands.

Applications:

- **Autonomous Driving:** An agent learns how to navigate traffic, brake, and accelerate to reach a destination safely.
- **Trading & Finance:** An agent learns to execute trades to maximize profit.
- **Supply Chain Management:** An agent learns to optimize inventory levels and logistics.

Important Notes / Key Takeaways:

- Learning is based on a reward system, not labeled data.
 - The agent learns by interacting with its environment.
 - The goal is to maximize cumulative reward.
 - Often used for decision-making problems.
-

5. Models of Machine Learning

Definition: Machine learning models can be broadly classified based on the underlying structure and assumptions they use to represent data and make predictions. These models fall into various categories, each suited for different types of problems and data. Understanding these models provides a fundamental framework for choosing the right approach for a given task.

5.1 Geometric Models

Definition: Geometric models are a class of machine learning models that are based on the concept of geometry in high-dimensional space. These models create decision boundaries, lines, or planes to separate or group data points based on their features. The spatial relationship between data points is key to how these models function.

Detailed Explanation:

- **Feature Space:** Data points are represented as vectors or points in a multi-dimensional space, where each dimension corresponds to a feature.
- **Decision Boundaries:** The model's learning is a process of finding the optimal geometric structure (e.g., a hyperplane) to separate different classes (for classification) or to represent the data's underlying structure (for regression).
- **Key Algorithms:**
 - **Support Vector Machines (SVMs):** Finds the hyperplane with the maximum margin between two classes.
 - **k-Nearest Neighbors (k-NN):** Classifies a data point based on the majority class of its nearest neighbors in the feature space.
 - **Linear Regression:** Fits a straight line or plane to the data to predict a continuous value.

Mathematical / Technical Aspect: A linear classifier's decision boundary can be represented by a linear equation: $w \cdot x + b = 0$ Where:

- w is the weight vector (normal to the hyperplane)
- x is the input feature vector
- b is the bias (offset from the origin)

Real-World Examples:

1. **Spam Classification with SVM:** An SVM model can be used to separate "spam" and "not spam" emails by finding the optimal hyperplane in a multi-dimensional feature space where each dimension represents a word's frequency.
2. **Handwriting Recognition with k-NN:** An image of a handwritten digit is compared to a large database of labeled digits. The model identifies the k-nearest neighbors and predicts the digit based on the most common class among them.
3. **Predicting Stock Price with Linear Regression:** A simple geometric model fits a line to historical stock prices to predict future price trends.

Applications:

- **Computer Vision:** Object detection and image classification.
- **Bioinformatics:** Protein structure prediction.
- **Medical Diagnosis:** Classifying a tumor as benign or malignant.

Important Notes / Key Takeaways:

- Based on spatial relationships of data in feature space.
- Uses lines, planes, and other geometric shapes to model data.
- Includes algorithms like SVM, k-NN, and linear regression.

5.2 Probabilistic Models

Definition: Probabilistic models are a class of machine learning models that use the principles of probability theory to model the relationships within the data. These models assume that the data is generated by an underlying probabilistic process and aim to find the most probable explanation for the observed data. They provide not just a prediction but also a measure of confidence or uncertainty.

Detailed Explanation:

- **Learning from Probability Distributions:** These models treat the relationships between variables as a matter of probability. They learn the probability distributions of the data and use these to make predictions.
- **Uncertainty Quantification:** A key advantage is the ability to provide a confidence score along with the prediction, which is crucial in many high-stakes applications like medical diagnosis or financial trading.
- **Key Algorithms:**
 - **Naive Bayes:** Based on Bayes' theorem, it assumes that features are conditionally independent.
 - **Bayesian Networks:** Represents relationships between variables as a directed acyclic graph.
 - **Probabilistic Latent Semantic Analysis (PLSA):** Used for topic modeling in text documents.

Mathematical / Technical Aspect: Naive Bayes uses Bayes' theorem to calculate the probability of a class given a set of features: $P(A|B) = [P(B|A) * P(A)] / P(B)$ Where:

- $P(A|B)$ is the posterior probability of class A given feature B.
- $P(B|A)$ is the likelihood of feature B given class A.
- $P(A)$ is the prior probability of class A.
- $P(B)$ is the prior probability of feature B.

Real-World Examples:

1. **Email Spam Detection with Naive Bayes:** The model calculates the probability that an email is spam based on the presence of certain keywords (e.g., "free," "money," "prize").
2. **Medical Diagnosis with Bayesian Networks:** A network can model the relationships between symptoms, test results, and diseases to calculate the probability of a specific disease given a patient's data.
3. **Language Translation:** Probabilistic models are used in some machine translation systems to find the most probable translation for a given sentence.

Applications:

- **Medical Informatics:** Disease diagnosis and risk assessment.
- **Spam Filtering:** Identifying junk emails.

- **Natural Language Processing:** Text classification and topic modeling.

Important Notes / Key Takeaways:

- Based on probability and statistics.
- Provides confidence scores or uncertainty measures with predictions.
- Includes Naive Bayes and Bayesian Networks.

5.3 Logical Models

Definition: Logical models are a class of machine learning models that use logical rules and decision-making structures to make predictions. Unlike geometric models that use continuous values and spatial relationships, logical models use a series of discrete, human-readable "if-then" rules. This makes their decision-making process highly interpretable and transparent.

Detailed Explanation:

- **Rule-Based Systems:** The model learns a set of rules from the data, which can be easily understood and verified by a human.
- **Hierarchical Decision-Making:** These models often build a tree-like structure where each node represents a decision based on a feature, leading to a final outcome.
- **Key Algorithms:**
 - **Decision Trees:** A tree-like model of decisions and their possible consequences. Each internal node represents a test on an attribute, each branch represents the outcome of the test, and each leaf node represents a class label.
 - **Rule-Based Classifiers:** A system that uses a set of "if-then" rules to classify data.
 - **Random Forests:** An ensemble method that builds multiple decision trees and combines their predictions.

Real-World Examples:

1. **Loan Approval Decision Tree:** A bank uses a decision tree to approve or deny a loan. The tree might have rules like "IF income > \$50k AND credit score > 700 THEN approve loan."

2. **Customer Churn Prediction:** A logical model can predict which customers are likely to cancel their subscription with rules like "IF customer has not logged in for 30 days AND has made less than 5 purchases THEN predict churn."
3. **Medical Diagnosis:** A doctor's decision-making process can be modeled as a decision tree, where each step (e.g., checking a symptom) leads to a different path and diagnosis.

Applications:

- **Expert Systems:** Emulating the decision-making of a human expert.
- **Business Intelligence:** Creating clear, actionable rules for business processes.
- **Credit Scoring:** Generating transparent, rule-based systems for risk assessment.

Important Notes / Key Takeaways:

- Uses a set of "if-then" rules.
- Highly interpretable and easy to understand.
- Includes decision trees and rule-based classifiers.

5.4 Grouping and Grading Models

Definition: Grouping and grading models are a broad class that can be a combination of other model types. **Grouping models** (often synonymous with clustering) focus on dividing a dataset into meaningful groups based on similarity. **Grading models** (often synonymous with regression) focus on predicting a numerical grade or score for a data point.

Detailed Explanation:

- **Grouping (Clustering):**
 - **Purpose:** To find and categorize data points into groups (clusters) where data points within a group are more similar to each other than to those in other groups.
 - **Model Type:** Unsupervised learning.
 - **Examples:** K-Means, Hierarchical Clustering.
- **Grading (Regression):**
 - **Purpose:** To predict a continuous output value or a "grade" for a given input.
 - **Model Type:** Supervised learning.
 - **Examples:** Linear Regression, Ridge Regression, Lasso Regression.

- **Hybrid Nature:** Many algorithms can be used for both grouping and grading, depending on how they are applied. For example, a model might first group data and then apply a grading model within each group.

Real-World Examples:

1. **Customer Segmentation (Grouping):** An e-commerce company groups its customers into segments like "loyal shoppers," "new visitors," and "seasonal buyers" to target them with specific promotions.
2. **Credit Score Prediction (Grading):** A bank uses a grading model to predict a customer's credit score (a numerical value) based on their financial history.
3. **Product Recommendation (Grouping & Grading):** A streaming service might first group users with similar viewing habits (grouping) and then, for a new user, grade the likelihood of them liking a particular movie based on what others in their group watched (grading).

Applications:

- **Marketing:** Customer segmentation.
- **Finance:** Credit risk scoring and stock price prediction.
- **Healthcare:** Patient risk stratification.

Important Notes / Key Takeaways:

- Grouping is about finding clusters or segments (unsupervised).
- Grading is about predicting a continuous numerical value (supervised).
- These concepts are often applied together in complex systems.

5.5 Parametric and Non-Parametric Models

Definition: This is a classification based on the model's assumptions about the underlying data distribution. **Parametric models** make strong assumptions about the form of the data's distribution. **Non-parametric models**, on the other hand, make no or very few assumptions about the data's distribution.

Detailed Explanation:

- **Parametric Models:**

- **Assumption:** Assume the data follows a specific, predefined functional form (e.g., a straight line, a bell curve).
- **Learning:** The learning process involves finding a fixed number of parameters (coefficients) for that function.
- **Advantages:** Simpler, faster to train, and require less data.
- **Disadvantages:** Can be inaccurate if the data does not fit the assumed form.
- **Examples:** Linear Regression, Logistic Regression, Naive Bayes.
- **Non-Parametric Models:**
 - **Assumption:** Make no assumptions about the data distribution. The model's structure is determined by the data itself.
 - **Learning:** The model's complexity grows with the amount of data.
 - **Advantages:** More flexible, can fit a wide range of data patterns.
 - **Disadvantages:** Require more data, can be slower to train, and are more prone to overfitting.
 - **Examples:** Decision Trees, k-Nearest Neighbors, Support Vector Machines (with certain kernels).

Real-World Examples:

1. Predicting Sales:

- **Parametric:** Using a linear regression model assumes that sales are a linear function of advertising spend. You just need to find the two parameters (slope and intercept).
- **Non-Parametric:** A decision tree model would learn a complex set of rules from the sales data, without assuming a linear relationship.

2. Image Classification:

- **Parametric:** A simple logistic regression model would assume a linear relationship between image pixel values and the class.
- **Non-Parametric:** A k-NN model makes no such assumption; it simply classifies a new image based on its similarity to existing ones in the dataset.

3. Customer Lifetime Value (CLV) Prediction:

- **Parametric:** You might assume CLV is a linear function of a customer's age and past purchases.
- **Non-Parametric:** A random forest model would learn complex, non-linear relationships to predict CLV without a predefined form.

Applications:

- **Parametric:** Used when the relationship between variables is simple and well-understood.
- **Non-Parametric:** Used when the relationship is complex, and there is a lot of data available.

Important Notes / Key Takeaways:

- Parametric models have a fixed number of parameters and assume a data distribution.
- Non-parametric models do not assume a distribution and their complexity grows with the data.
- The choice between them depends on the data size and complexity of the problem.

Summary Table

Main Topic Name	Brief Definition	3 Real-World Examples
Introduction to Machine Learning	A subfield of AI that enables computers to learn from data without explicit programming. It uses algorithms and statistical models to identify patterns and make predictions, with a focus on automatic improvement over time.	1. Email Spam Detection 2. Netflix Recommendations 3. Predictive Text on Smartphones
ML vs Traditional Programming	Traditional programming requires a developer to write explicit rules for a computer, while ML provides data and an algorithm, allowing the computer to learn the rules on its own from the data.	1. Weather Prediction 2. Face Recognition 3. Language Translation
ML vs AI vs Data Science	AI is the broad concept of creating intelligent machines. ML is a specific method used to achieve AI by learning from data. Data Science is an interdisciplinary field that uses various techniques, including ML, to extract insights from data.	1. Self-Driving Cars 2. Fraud Detection in Banking 3. E-commerce Recommendation Systems
Supervised Learning	A type of learning where an algorithm is trained on a labeled dataset, learning a mapping from input features to known output values to make future predictions.	1. House Price Prediction 2. Credit Risk Assessment 3. Image Classification

Unsupervised Learning	A type of learning that uses an unlabeled dataset to find hidden patterns and structures within the data on its own, without any prior knowledge of the outcomes.	1. Customer Segmentation 2. Anomaly Detection 3. Market Basket Analysis
Semi-Supervised Learning	A hybrid approach that uses a small amount of labeled data and a large amount of unlabeled data to train a model, leveraging the unlabeled data to improve performance.	1. Web Content Classification 2. Speech Recognition 3. Medical Image Segmentation
Reinforcement Learning	A paradigm where an "agent" learns to take actions in an "environment" to maximize a cumulative "reward" through a process of trial and error and feedback.	1. AI for Game Playing (e.g., Chess) 2. Robotics (e.g., learning to walk) 3. Autonomous Driving
Geometric Models	Models that represent data in a multi-dimensional space and use geometric concepts like lines, planes, and distances to create decision boundaries or group data points.	1. Spam Classification (SVM) 2. Handwriting Recognition (k-NN) 3. Stock Price Prediction (Linear Regression)
Probabilistic Models	Models that use probability theory to represent relationships within the data, providing not only predictions but also a measure of confidence or uncertainty.	1. Email Spam Filtering (Naive Bayes) 2. Medical Diagnosis (Bayesian Networks) 3. Language Translation
Logical Models	Models that use a series of discrete "if-then" rules to make predictions, resulting in highly interpretable and transparent decision-making processes.	1. Loan Approval Decision Tree 2. Customer Churn Prediction 3. Medical Diagnosis
Grouping & Grading Models	Grouping models segment data into clusters based on similarity. Grading models predict a continuous, numerical value or score. Both	

Thank-You

We sincerely appreciate your time and attention in reviewing this document. Your feedback and support mean a lot to us, and we look forward to building something valuable together.



Visit site to access further Notes!!!

✨ Why Choose Us?

- 📖 Free & Updated Notes for All Subjects
- 🎯 Exam-Focused Content
- 🔍 Easy Navigation by Unit & Topic
- 💻 Works on Mobile & Desktop

Branches: **CS | IT | AI&DS | AI&ML**



Notes By: puneuniversity.netlify.app

🔗 Visit us: <https://cutt.ly/lrHFgvYl>